

```

# Import libraries
import torch
import torch.nn as nn
import torchvision
import torchvision.transforms as transforms
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
import time
from queue import PriorityQueue
from sklearn.manifold import TSNE
from sklearn.decomposition import PCA

# Device configuration
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
device

device(type='cuda', index=0)

# CIFAR10 dataset
whole_train_dataset = torchvision.datasets.CIFAR10(root="./data", train=True, transform=None,
                                                    download=True)


test_dataset = torchvision.datasets.CIFAR10(root="./data", train=False, transform=None)
train_dataset, valid_dataset = torch.utils.data.random_split(whole_train_dataset, [
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', '

Downloading https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz to ./da
100% 170498071/170498071 [00:13<00:00,
13555463 82it/s]

# Batch size of the dataset
batch_size = 64
# Data loader
train_loader = torch.utils.data.DataLoader(dataset=train_dataset, batch_size=batch_size,
                                             shuffle=True)
valid_loader = torch.utils.data.DataLoader(dataset=valid_dataset, batch_size=batch_size,
                                             shuffle=True)
test_loader = torch.utils.data.DataLoader(dataset=test_dataset, batch_size=batch_size,
                                           shuffle=True)

# Visualise a batch of images
plt.subplots(1,1,figsize=[20,20])
image, label = iter(train_loader).next()
grid_image = torchvision.utils.make_grid(image)
permuted_image = torch.permute(grid_image, (1,2,0))
plt.imshow(permuted_image)
plt.show()

```



```

# Original CNN
num_class = 10
class ConvNet(nn.Module):
    def __init__(self, num_class=10):
        super(ConvNet, self).__init__()
        self._layer1 = nn.Sequential(
            nn.Conv2d(in_channels=3, out_channels=32, kernel_size=3, stride=1, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2)
        ) # Output size = (32, 16, 16)
        self._layer2 = nn.Sequential(
            nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, stride=1, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2)
        ) # Output size = (64, 8, 8)
        self._layer3 = nn.Sequential(
            nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3, stride=1, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2)
        ) # Output size = (128, 4, 4)
        self._layer4 = nn.Sequential(
            nn.Conv2d(in_channels=128, out_channels=2048, kernel_size=4, stride=1, padding=1),
            nn.ReLU()
        ) # Output size = (2048, 1, 1)
        self._layer5 = nn.Linear(in_features=2048*1*1, out_features=4096)
        self._layer6 = nn.Linear(in_features=4096, out_features=num_class)

        self._layer1_out = None
        self._layer2_out = None
        self._layer3_out = None
        self._layer4_out = None
        self._layer5_out = None
        self._layer6_out = None

    def forward(self, x):
        self._layer1_out = self._layer1(x)
        self._layer2_out = self._layer2(self._layer1_out)
        self._layer3_out = self._layer3(self._layer2_out)
        self._layer4_out = self._layer4(self._layer3_out)
        self._layer5_out = self._layer5(self._layer4_out.reshape(self._layer4_out.size()[0]))
        self._layer6_out = self._layer6(self._layer5_out)
        return self._layer6_out

    def get_layer(self, num_layer):
        if num_layer == 1:
            return self._layer1_out
        elif num_layer == 2:
            return self._layer2_out
        elif num_layer == 3:
            return self._layer3_out
        elif num_layer == 4:
            return self._layer4_out

```

```

elif num_layer == 5:
    return self._layer5_out
else:
    return self._layer6_out

```

# CNN with 1 convolutional block

```
num_class = 10
```

```
class ConvNet1layer(nn.Module):
```

```

    def __init__(self, num_class=10):
        super(ConvNet1layer, self).__init__()
        self._layer1 = nn.Sequential(
            nn.Conv2d(in_channels=3, out_channels=32, kernel_size=3, stride=1, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2)
        ) # Output size = (32, 16, 16)
        self._layer2 = nn.Linear(in_features=32*16*16, out_features=4096)
        self._layer3 = nn.Linear(in_features=4096, out_features=num_class)

        self._layer1_out = None
        self._layer2_out = None
        self._layer3_out = None

```

```

    def forward(self, x):
        self._layer1_out = self._layer1(x)
        self._layer2_out = self._layer2(self._layer1_out.reshape(self._layer1_out.size()[0], -1))
        self._layer3_out = self._layer3(self._layer2_out)
        return self._layer3_out

```

```
    def get_layer(self, num_layer):
```

```

        if num_layer == 1:
            return self._layer1_out
        elif num_layer == 2:
            return self._layer2_out
        elif num_layer == 3:
            return self._layer3_out

```

# CNN with 2 convolutional block

```
num_class = 10
```

```
class ConvNet2layer(nn.Module):
```

```

    def __init__(self, num_class=10):
        super(ConvNet2layer, self).__init__()
        self._layer1 = nn.Sequential(
            nn.Conv2d(in_channels=3, out_channels=32, kernel_size=3, stride=1, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2)
        ) # Output size = (32, 16, 16)
        self._layer2 = nn.Sequential(
            nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, stride=1, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2)
        ) # Output size = (64, 8, 8)
        self._layer3 = nn.Linear(in_features=64*8*8, out_features=4096)
        self._layer4 = nn.Linear(in_features=4096, out_features=num_class)

```

```

        self._layer1_out = None
        self._layer2_out = None
        self._layer3_out = None
        self._layer4_out = None

    def forward(self, x):
        self._layer1_out = self._layer1(x)
        self._layer2_out = self._layer2(self._layer1_out)
        self._layer3_out = self._layer3(self._layer2_out.reshape(self._layer2_out.s
        self._layer4_out = self._layer4(self._layer3_out)
        return self._layer4_out

    def get_layer(self, num_layer):
        if num_layer == 1:
            return self._layer1_out
        elif num_layer == 2:
            return self._layer2_out
        elif num_layer == 3:
            return self._layer3_out
        elif num_layer == 4:
            return self._layer4_out

# CNN with 1 convolutional block
num_class = 10
class ConvNet3layer(nn.Module):
    def __init__(self, num_class=10):
        super(ConvNet3layer, self).__init__()
        self._layer1 = nn.Sequential(
            nn.Conv2d(in_channels=3, out_channels=32, kernel_size=3, stride=1, padd
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2)
        ) # Output size = (32, 16, 16)
        self._layer2 = nn.Sequential(
            nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, stride=1, pad
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2)
        ) # Output size = (64, 8, 8)
        self._layer3 = nn.Sequential(
            nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3, stride=1, pa
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2)
        ) # Output size = (128, 4, 4)
        self._layer4 = nn.Linear(in_features=128*4*4, out_features=4096)
        self._layer5 = nn.Linear(in_features=4096, out_features=num_class)

        self._layer1_out = None
        self._layer2_out = None
        self._layer3_out = None
        self._layer4_out = None
        self._layer5_out = None

    def forward(self, x):
        self._layer1_out = self._layer1(x)

```

```

self._layer2_out = self._layer2(self._layer1_out)
self._layer3_out = self._layer3(self._layer2_out)
self._layer4_out = self._layer4(self._layer3_out.reshape(self._layer3_out.s
self._layer5_out = self._layer5(self._layer4_out)
return self._layer5_out

```

```
def get_layer(self, num_layer):
```

```

    if num_layer == 1:
        return self._layer1_out
    elif num_layer == 2:
        return self._layer2_out
    elif num_layer == 3:
        return self._layer3_out
    elif num_layer == 4:
        return self._layer4_out
    elif num_layer == 5:
        return self._layer5_out

```

```
# CNN with no max pooling layers
```

```
num_class = 10
```

```
class ConvNetNoMax(nn.Module):
```

```
    def __init__(self, num_class=10):
```

```
        super(ConvNetNoMax, self).__init__()
```

```
        self._layer1 = nn.Sequential(
```

```
            nn.Conv2d(in_channels=3, out_channels=32, kernel_size=3, stride=1, padd
            nn.ReLU())
```

```
        ) # Output size = (32, 32, 32)
```

```
        self._layer2 = nn.Sequential(
```

```
            nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, stride=1, pad
            nn.ReLU())
```

```
        ) # Output size = (64, 32, 32)
```

```
        self._layer3 = nn.Linear(in_features=64*32*32, out_features=4096)
```

```
        self._layer4 = nn.Linear(in_features=4096, out_features=num_class)
```

```
        self._layer1_out = None
```

```
        self._layer2_out = None
```

```
        self._layer3_out = None
```

```
        self._layer4_out = None
```

```
    def forward(self, x):
```

```
        self._layer1_out = self._layer1(x)
```

```
        self._layer2_out = self._layer2(self._layer1_out)
```

```
        self._layer3_out = self._layer3(self._layer2_out.reshape(self._layer2_out.s
```

```
        self._layer4_out = self._layer4(self._layer3_out)
```

```
        return self._layer4_out
```

```
    def get_layer(self, num_layer):
```

```
        if num_layer == 1:
```

```
            return self._layer1_out
```

```
        elif num_layer == 2:
```

```
            return self._layer2_out
```

```
        elif num_layer == 3:
```

```
            return self._layer3_out
```

```

        elif num_layer == 4:
            return self._layer4_out

# DNN
num_class = 10
class DenseNet(nn.Module):
    def __init__(self, num_class=10):
        super(DenseNet, self).__init__()

        self._layer1 = nn.Sequential(
            nn.Linear(in_features=32*32*3, out_features=512),
            nn.ReLU()
        ) # Output size = (1, 512)
        self._layer2 = nn.Sequential(
            nn.Linear(in_features=512, out_features=1024),
            nn.ReLU()
        ) # Output size = (1, 1024)
        self._layer3 = nn.Sequential(
            nn.Linear(in_features=1024, out_features=2048),
            nn.ReLU()
        ) # Output size = (1, 2048)
        self._layer4 = nn.Sequential(
            nn.Linear(in_features=2048, out_features=4096),
            nn.ReLU()
        ) # Output size = (1, 4096)
        self._layer5 = nn.Linear(in_features=4096, out_features=num_class)

        self._layer1_out = None
        self._layer2_out = None
        self._layer3_out = None
        self._layer4_out = None
        self._layer5_out = None

    def forward(self, x):
        x = x.reshape(-1, 32*32*3)
        self._layer1_out = self._layer1(x)
        self._layer2_out = self._layer2(self._layer1_out)
        self._layer3_out = self._layer3(self._layer2_out)
        self._layer4_out = self._layer4(self._layer3_out)
        self._layer5_out = self._layer5(self._layer4_out)
        return self._layer5_out

    def get_layer(self, num_layer):
        if num_layer == 1:
            return self._layer1_out
        elif num_layer == 2:
            return self._layer2_out
        elif num_layer == 3:
            return self._layer3_out
        elif num_layer == 4:
            return self._layer4_out
        else:
            return self._layer5_out

```

```

# Train function for neural networks
def train(model, criterion, learning_rate, optimizer, train_loader, valid_loader, n
    # Parameters to be defined
    total_steps = len(train_loader)
    train_accuracy = 0
    train_accuracy_list = []
    valid_accuracy = 0
    best_valid_accuracy = 0
    no_improvement = 0
    valid_accuracy_list = []
    loss_list = []
    # Forward feed and Optimization
    for epoch in range(num_epoch):
        for i, (image, label) in enumerate(train_loader):
            # Load a batch of data
            image = image.to(device)
            label = label.to(device)

            # Put data into the neural network
            output = model(image)
            # Evaluate the loss
            loss = criterion(output, label)
            # Backpropagation and optimization
            optimizer.zero_grad()
            loss.backward()
            optimizer.step()
            # Save loss into a list for the learning curve
            loss_list.append(loss.item())

            # Neural network evaluation in the training set
            model.eval()
            with torch.no_grad():
                train_correct = 0
                train_total = 0
                for image, label in train_loader:
                    image = image.to(device)
                    label = label.to(device)

                    output = model(image)
                    _, predicted = torch.max(output.data, 1)
                    train_total += label.size(0)
                    train_correct += (predicted == label).sum().item()

            # Save training set accuracy for the learning curve
            train_accuracy = train_correct/train_total*100
            train_accuracy_list.append(train_accuracy)

            # Neural network evaluation in the validation set
            with torch.no_grad():
                valid_correct = 0
                valid_total = 0
                for image, label in valid_loader:
                    image = image.to(device)
                    label = label.to(device)

```



```

        output = model(image)
        _, predicted = torch.max(output.data, 1)
        valid_total += label.size(0)
        valid_correct += (predicted == label).sum().item()

    # Save training set accuracy for the learning curve
    valid_accuracy = valid_correct/valid_total*100
    valid_accuracy_list.append(valid_accuracy)
    print(f"Epoch[{epoch+1}/{num_epoch}], Loss:{loss.item()}, Training Accuracy:


    # Early stopping
    if valid_accuracy > best_valid_accuracy:
        best_valid_accuracy = valid_accuracy
        no_improvement = 0
    else:
        no_improvement += 1
    if no_improvement == n_no_improvement:
        break
    return model, loss_list, train_accuracy_list, valid_accuracy_list

# Learning curve function
def learning_curve(save=False, path=None, title=None, train_accuracy_list=None, val
    fig, ax = plt.subplots(1,1,figsize=[8,8])
    ax.plot([i for i in range(len(train_accuracy_list))], train_accuracy_list, mark
    ax.plot([i for i in range(len(valid_accuracy_list))], valid_accuracy_list, mark
    ax.plot([i for i in range(len(loss_list))], loss_list, marker=".", label="train
    ax.set_title(title)
    ax.set_xlabel("Epoch")
    ax.set_ylabel("Accuracy(%)")
    plt.legend()
    if save == True:
        fig.savefig(path)

# Train an original CNN
cnn_model = ConvNet(num_class=10).to(device)
criterion = nn.CrossEntropyLoss()
learning_rate = 0.1
optimizer = torch.optim.SGD(params=cnn_model.parameters(), lr=learning_rate)
train_loader = train_loader
valid_loader = valid_loader
num_epoch = 30
n_no_improvement = 3
cnn_model, loss_list, train_accuracy_list, valid_accuracy_list = train(cnn_model, c
learning_curve(save=True, path="cnn_model_original_lc", train_accuracy_list=train_a

```

```
Epoch[1/30], Loss:1.834172248840332, Training Accuracy: 39.9075%, Validation Accuracy: 39.9075%
Epoch[2/30], Loss:1.1700890064239502, Training Accuracy: 51.370000000000005, Validation Accuracy: 51.370000000000005
Epoch[3/30], Loss:1.1501590013504028, Training Accuracy: 59.545%, Validation Accuracy: 59.545%
Epoch[4/30], Loss:1.097917914390564, Training Accuracy: 61.4425%, Validation Accuracy: 61.4425%
Epoch[5/30], Loss:0.8922512531280518, Training Accuracy: 74.21%, Validation Accuracy: 74.21%
Epoch[6/30], Loss:0.5716832876205444, Training Accuracy: 79.535%, Validation Accuracy: 79.535%
Epoch[7/30], Loss:0.49166393280029297, Training Accuracy: 87.88%, Validation Accuracy: 87.88%
Epoch[8/30], Loss:0.6343632936477661, Training Accuracy: 91.435%, Validation Accuracy: 91.435%
Epoch[9/30], Loss:0.3750077486038208, Training Accuracy: 92.8425%, Validation Accuracy: 92.8425%
Epoch[10/30], Loss:0.22277632355690002, Training Accuracy: 96.2525%, Validation Accuracy: 96.2525%
```





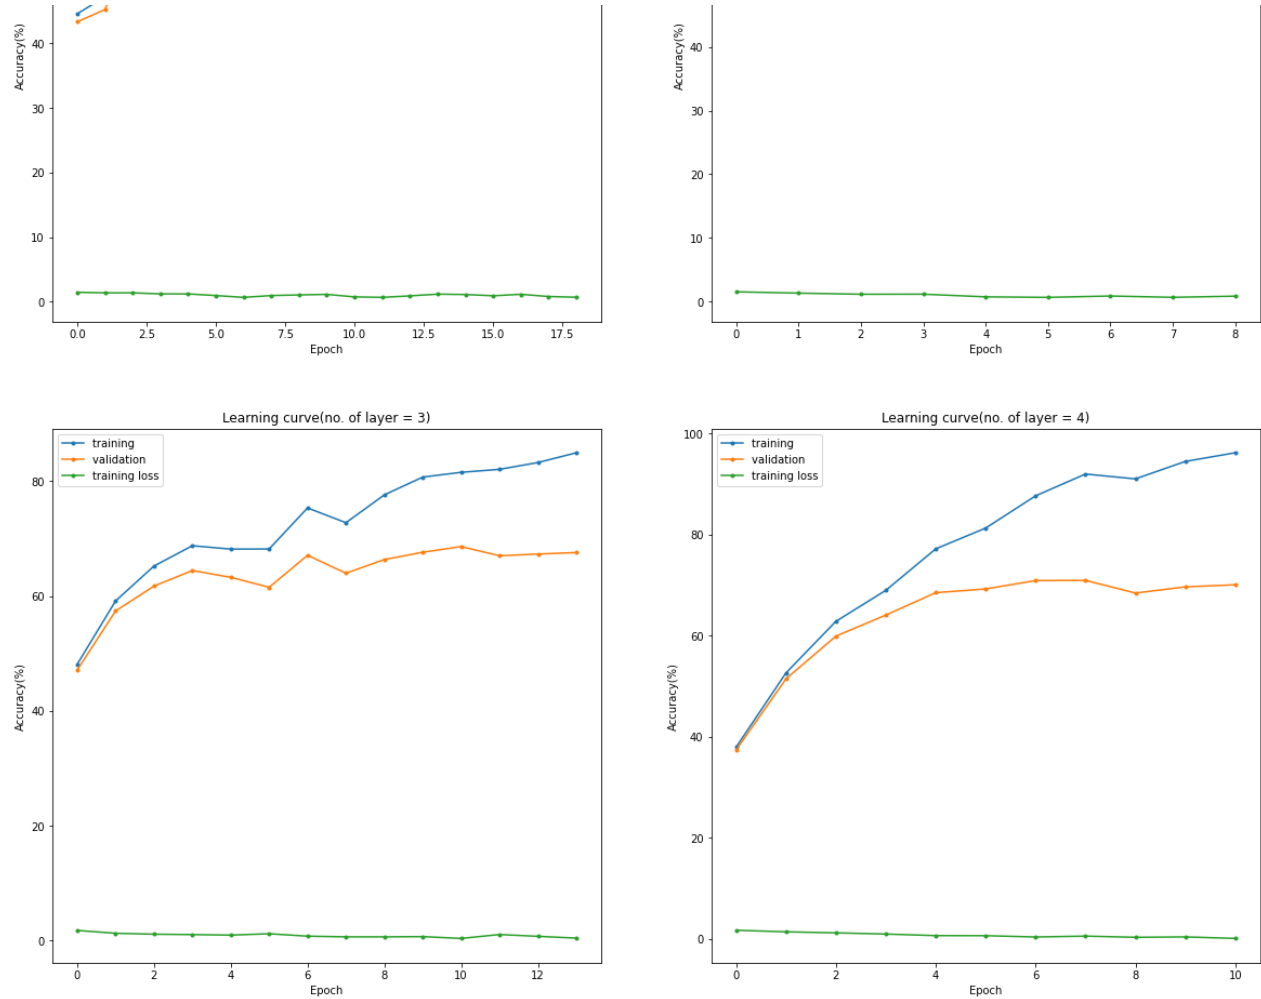
```
# Train CNNs with 1, 2, 3, 4 convolutional blocks
fig, ax = plt.subplots(2,2,figsize=[20,20])
axis = [(0,0),(0,1),(1,0),(1,1)]
train_accuracy_max = []
valid_accuracy_max = []
loss_min = []
for index, CNN in enumerate([ConvNet1layer, ConvNet2layer, ConvNet2layer, ConvNet])
    cnn_model = CNN(num_class=10).to(device)
    criterion = nn.CrossEntropyLoss()
    learning_rate = 0.1
    optimizer = torch.optim.SGD(params=cnn_model.parameters(), lr=learning_rate)
    train_loader = train_loader
    valid_loader = valid_loader
    num_epoch = 30
    n_no_improvement = 3
    cnn_model, loss_list, train_accuracy_list, valid_accuracy_list = train(cnn_model,
    train_accuracy_max.append(max(train_accuracy_list))
    valid_accuracy_max.append(max(valid_accuracy_list))
    loss_min.append(min(loss_list))
    row, column = axis[index]
    ax[row][column].plot([i for i in range(len(train_accuracy_list))], train_accuracy
    ax[row][column].plot([i for i in range(len(valid_accuracy_list))], valid_accuracy
    ax[row][column].plot([i for i in range(len(loss_list))], loss_list, marker=".", l
    ax[row][column].set_title(f"Learning curve(no. of layer = {index+1})")
    ax[row][column].set_xlabel("Epoch")
    ax[row][column].set_ylabel("Accuracy(%)")
    ax[row][column].legend()
```

```


Epoch[1/30], Loss:1.4435992240905762, Training Accuracy: 44.5575%, Validation
Epoch[2/30], Loss:1.36066472530365, Training Accuracy: 47.2025%, Validation
Epoch[3/30], Loss:1.3569222688674927, Training Accuracy: 59.99250000000001%
Epoch[4/30], Loss:1.1947115659713745, Training Accuracy: 60.614999999999995
Epoch[5/30], Loss:1.1874756813049316, Training Accuracy: 62.517500000000005
Epoch[6/30], Loss:0.9373486042022705, Training Accuracy: 66.1225%, Validation
Epoch[7/30], Loss:0.6569656729698181, Training Accuracy: 64.6525%, Validation
Epoch[8/30], Loss:0.9363547563552856, Training Accuracy: 67.755%, Validation
Epoch[9/30], Loss:1.0311768054962158, Training Accuracy: 70.4025%, Validation
Epoch[10/30], Loss:1.120340347290039, Training Accuracy: 68.07%, Validation
Epoch[11/30], Loss:0.7297576069831848, Training Accuracy: 72.1975%, Validation
Epoch[12/30], Loss:0.6615102291107178, Training Accuracy: 73.1325%, Validation
Epoch[13/30], Loss:0.8888604044914246, Training Accuracy: 72.3725%, Validation
Epoch[14/30], Loss:1.1530753374099731, Training Accuracy: 73.69250000000001
Epoch[15/30], Loss:1.0986626148223877, Training Accuracy: 73.4825%, Validation
Epoch[16/30], Loss:0.9024112820625305, Training Accuracy: 75.53%, Validation
Epoch[17/30], Loss:1.1285609006881714, Training Accuracy: 73.0375%, Validation
Epoch[18/30], Loss:0.7884595990180969, Training Accuracy: 75.8725%, Validation
Epoch[19/30], Loss:0.6789387464523315, Training Accuracy: 71.6275%, Validation
Epoch[1/30], Loss:1.5265092849731445, Training Accuracy: 51.5375%, Validation
Epoch[2/30], Loss:1.3237429857254028, Training Accuracy: 61.482499999999995
Epoch[3/30], Loss:1.1452070474624634, Training Accuracy: 63.77%, Validation
Epoch[4/30], Loss:1.154433250427246, Training Accuracy: 68.1875%, Validation
Epoch[5/30], Loss:0.737276554107666, Training Accuracy: 71.9725%, Validation
Epoch[6/30], Loss:0.6562880277633667, Training Accuracy: 73.6575%, Validation
Epoch[7/30], Loss:0.869956374168396, Training Accuracy: 73.775%, Validation
Epoch[8/30], Loss:0.6633521318435669, Training Accuracy: 75.0325%, Validation
Epoch[9/30], Loss:0.845159649848938, Training Accuracy: 76.9875%, Validation
Epoch[1/30], Loss:1.732466697692871, Training Accuracy: 48.089999999999996%
Epoch[2/30], Loss:1.2401726245880127, Training Accuracy: 59.137499999999996
Epoch[3/30], Loss:1.074942708015442, Training Accuracy: 65.225%, Validation
Epoch[4/30], Loss:1.0060639381408691, Training Accuracy: 68.77%, Validation
Epoch[5/30], Loss:0.914389967918396, Training Accuracy: 68.1825%, Validation
Epoch[6/30], Loss:1.1533972024917603, Training Accuracy: 68.2%, Validation
Epoch[7/30], Loss:0.7320619821548462, Training Accuracy: 75.3525%, Validation
Epoch[8/30], Loss:0.613974928855896, Training Accuracy: 72.7825%, Validation
Epoch[9/30], Loss:0.6109513640403748, Training Accuracy: 77.63250000000001%
Epoch[10/30], Loss:0.6543534994125366, Training Accuracy: 80.72749999999999
Epoch[11/30], Loss:0.33957940340042114, Training Accuracy: 81.5625%, Validation
Epoch[12/30], Loss:1.0179082155227661, Training Accuracy: 82.0825%, Validation
Epoch[13/30], Loss:0.7025007009506226, Training Accuracy: 83.265%, Validation
Epoch[14/30], Loss:0.394182026386261, Training Accuracy: 84.9675%, Validation
Epoch[1/30], Loss:1.7090626955032349, Training Accuracy: 37.9%, Validation
Epoch[2/30], Loss:1.396533489227295, Training Accuracy: 52.5975%, Validation
Epoch[3/30], Loss:1.1965198516845703, Training Accuracy: 62.7925%, Validation
Epoch[4/30], Loss:0.9559071660041809, Training Accuracy: 68.92%, Validation
Epoch[5/30], Loss:0.6474428176879883, Training Accuracy: 77.125%, Validation
Epoch[6/30], Loss:0.6329729557037354, Training Accuracy: 81.2325%, Validation
Epoch[7/30], Loss:0.3778480291366577, Training Accuracy: 87.605%, Validation
Epoch[8/30], Loss:0.5472712516784668, Training Accuracy: 91.92%, Validation
Epoch[9/30], Loss:0.3155595660209656, Training Accuracy: 90.96%, Validation
Epoch[10/30], Loss:0.3944661617279053, Training Accuracy: 94.415%, Validation
Epoch[11/30], Loss:0.09311220794916153, Training Accuracy: 96.1225%, Validation

```





```
# Plot the maximum accuracy in training set, validation set and minimum loss of CNN
fig, ax = plt.subplots(1,2,figsize=[15,5])
ax[0].plot([i+1 for i in range(len(train_accuracy_max))], train_accuracy_max, marker="o")
ax[0].plot([i+1 for i in range(len(valid_accuracy_max))], valid_accuracy_max, marker="o")
ax[1].plot([i+1 for i in range(len(loss_min))], loss_min, marker=".", label="training loss")
ax[0].set_title(f"Maximum accuracy in N-layer")
ax[0].set_xlabel("Epoch")
ax[0].set_ylabel("Accuracy(%)")
ax[0].legend()
ax[1].set_title(f"Minimum loss in N-layer")
ax[1].set_xlabel("Epoch")
ax[1].set_ylabel("Loss")
ax[1].legend()
fig.savefig("cnn_layer_summary.png")
```



```
# Train CNNs with 10^-5 to 10^0 L2 regularization
fig, ax = plt.subplots(2,3, figsize=[15,15])
row_column = [(0,0), (0,1), (0,2), (1,0), (1,1), (1,2)]
index = 0
for index, alpha in enumerate([-5, -4, -3, -2, -1, 0]):
    cnn_model = ConvNet(num_class=10).to(device)
    criterion = nn.CrossEntropyLoss()
    learning_rate = 0.1
    optimizer = torch.optim.SGD(params=cnn_model.parameters(), lr=learning_rate, weight_decay=10**alpha)
    train_loader = train_loader
    valid_loader = valid_loader
    num_epoch = 50
    n_no_improvement = 5
    cnn_model, loss_list, train_accuracy_list, valid_accuracy_list = train(cnn_model,
    row, column = row_column[index])
    ax[row][column].plot([i for i in range(len(train_accuracy_list))], train_accuracy_list, marker="o")
    ax[row][column].plot([i for i in range(len(valid_accuracy_list))], valid_accuracy_list, marker="o")
    ax[row][column].plot([i for i in range(len(loss_list))], loss_list, marker=".", label="training loss")
    ax[row][column].set_title(f"Learning curve(L2:10^{alpha})")
    ax[row][column].set_xlabel("Epoch")
```

```
ax[row][column].set_ylabel("Accuracy(%)")  
ax[row][column].legend()
```

```
Epoch[1/50], Loss:1.7979432344436646, Training Accuracy: 37.0%, Validation
Epoch[2/50], Loss:1.31745183467865, Training Accuracy: 51.127500000000005%,
Epoch[3/50], Loss:1.158024549484253, Training Accuracy: 60.540000000000006%
Epoch[4/50], Loss:1.2576788663864136, Training Accuracy: 64.7175%, Validati
Epoch[5/50], Loss:0.8490377068519592, Training Accuracy: 73.165%, Validatio
Epoch[6/50], Loss:0.7441189885139465, Training Accuracy: 80.825%, Validatio
Epoch[7/50], Loss:0.5174120664596558, Training Accuracy: 86.1275%, Validati
Epoch[8/50], Loss:0.5251293778419495, Training Accuracy: 90.3725%, Validati
Epoch[9/50], Loss:0.39576414227485657, Training Accuracy: 94.1575%, Validat
Epoch[10/50], Loss:0.38876891136169434, Training Accuracy: 95.705%, Validat
Epoch[11/50], Loss:0.12799574434757233, Training Accuracy: 95.6%, Validatio
Epoch[12/50], Loss:0.30046820640563965, Training Accuracy: 97.1475%, Validat
Epoch[13/50], Loss:0.050701528787612915, Training Accuracy: 98.7025%, Valid
Epoch[14/50], Loss:0.09228236228227615, Training Accuracy: 98.105%, Validat
Epoch[15/50], Loss:0.040415484458208084, Training Accuracy: 98.507500000000
Epoch[16/50], Loss:0.05204831808805466, Training Accuracy: 98.3275%, Validat
Epoch[17/50], Loss:0.004636985715478659, Training Accuracy: 98.304999999999
Epoch[18/50], Loss:0.044425494968891144, Training Accuracy: 97.6%, Validati
Epoch[1/50], Loss:1.7394925355911255, Training Accuracy: 40.527499999999996
Epoch[2/50], Loss:1.2872847318649292, Training Accuracy: 54.1025%, Validati
Epoch[3/50], Loss:1.3444874286651611, Training Accuracy: 62.822500000000005
Epoch[4/50], Loss:0.8954426050186157, Training Accuracy: 70.21%, Validation
Epoch[5/50], Loss:1.1714483499526978, Training Accuracy: 73.33%, Validation
Epoch[6/50], Loss:0.5379042029380798, Training Accuracy: 81.38250000000001%
Epoch[7/50], Loss:0.4790600538253784, Training Accuracy: 85.66749999999999%
Epoch[8/50], Loss:0.7279185652732849, Training Accuracy: 86.13%, Validation
Epoch[9/50], Loss:0.34353113174438477, Training Accuracy: 93.63%, Validatio
Epoch[10/50], Loss:0.25099027156829834, Training Accuracy: 96.1425%, Validat
Epoch[11/50], Loss:0.1391584575176239, Training Accuracy: 96.76%, Validatio
Epoch[12/50], Loss:0.1813632696866989, Training Accuracy: 94.2975%, Validat
Epoch[13/50], Loss:0.10117306560277939, Training Accuracy: 97.08%, Validati
Epoch[14/50], Loss:0.15188603103160858, Training Accuracy: 98.352499999999
Epoch[1/50], Loss:1.7536437511444092, Training Accuracy: 38.0525%, Validati
Epoch[2/50], Loss:1.5623881816864014, Training Accuracy: 48.52%, Validation
Epoch[3/50], Loss:1.3785285949707031, Training Accuracy: 59.565%, Validatio
Epoch[4/50], Loss:0.9560930132865906, Training Accuracy: 67.09%, Validation
Epoch[5/50], Loss:0.8478828072547913, Training Accuracy: 70.635%, Validatio
Epoch[6/50], Loss:0.9023010730743408, Training Accuracy: 72.795%, Validatio
Epoch[7/50], Loss:0.9888664484024048, Training Accuracy: 81.765%, Validatio
Epoch[8/50], Loss:0.646549642086029, Training Accuracy: 84.125%, Validation
Epoch[9/50], Loss:0.5904465317726135, Training Accuracy: 85.065%, Validatio
Epoch[10/50], Loss:0.500035285949707, Training Accuracy: 92.335%, Validatio
Epoch[11/50], Loss:0.29693758487701416, Training Accuracy: 94.6875%, Validat
Epoch[12/50], Loss:0.12369435280561447, Training Accuracy: 96.6125%, Validat
Epoch[13/50], Loss:0.18576504290103912, Training Accuracy: 97.0425%, Validat
Epoch[14/50], Loss:0.2875842750072479, Training Accuracy: 93.102499999999
Epoch[15/50], Loss:0.12597142159938812, Training Accuracy: 97.5675%, Validat
Epoch[16/50], Loss:0.07494053989648819, Training Accuracy: 97.7225%, Validat
Epoch[1/50], Loss:1.8838781118392944, Training Accuracy: 27.665%, Validatio
Epoch[2/50], Loss:1.6017838716506958, Training Accuracy: 38.57%, Validation
Epoch[3/50], Loss:1.831206202507019, Training Accuracy: 42.4975%, Validatio
Epoch[4/50], Loss:1.821392297744751, Training Accuracy: 43.64%, Validation
Epoch[5/50], Loss:1.3605544567108154, Training Accuracy: 46.612500000000004
Epoch[6/50], Loss:1.265623688697815, Training Accuracy: 49.5325%, Validatio
Epoch[7/50], Loss:1.4581109285354614, Training Accuracy: 51.2625%, Validati
Epoch[8/50], Loss:1.4398549795150757, Training Accuracy: 53.0725%, Validati
Epoch[9/50], Loss:1.205130696296692, Training Accuracy: 56.9975%, Validatio
Epoch[10/50], Loss:1.4654110670089722, Training Accuracy: 58.194999999999
Epoch[11/50], Loss:1.29230797290802, Training Accuracy: 55.925000000000004%
Epoch[12/50], Loss:1.4419119358062744, Training Accuracy: 57.4225%, Validat
```

```

Epoch[13/50], Loss:1.3823338747024536, Training Accuracy: 60.785%, Validati
Epoch[14/50], Loss:1.020607352256775, Training Accuracy: 61.845000000000006
Epoch[15/50], Loss:0.9119833707809448, Training Accuracy: 63.295%, Validati
Epoch[16/50], Loss:1.155288577079773, Training Accuracy: 65.9375%, Validati
Epoch[17/50], Loss:0.9489369988441467, Training Accuracy: 67.32249999999999
Epoch[18/50], Loss:0.8850423097610474, Training Accuracy: 63.34999999999999
Epoch[19/50], Loss:1.0125056505203247, Training Accuracy: 64.022500000000001
Epoch[20/50], Loss:1.0995625257492065, Training Accuracy: 65.59%, Validatio
Epoch[21/50], Loss:1.2184432744979858, Training Accuracy: 61.5725%, Validat
Epoch[22/50], Loss:1.0323705673217773, Training Accuracy: 59.89499999999999
Epoch[1/50], Loss:2.2986488342285156, Training Accuracy: 9.945%, Validation
Epoch[2/50], Loss:2.3053767681121826, Training Accuracy: 9.950000000000001%
Epoch[3/50], Loss:2.3015899658203125, Training Accuracy: 9.950000000000001%
Epoch[4/50], Loss:2.3033056259155273, Training Accuracy: 10.127500000000001
Epoch[5/50], Loss:2.300581455230713, Training Accuracy: 10.0025%, Validatio
Epoch[6/50], Loss:2.3043503761291504, Training Accuracy: 10.0025%, Validati
Epoch[1/50], Loss:2.3018782138824463, Training Accuracy: 9.950000000000001%
Epoch[2/50], Loss:2.3026883602142334, Training Accuracy: 10.055%, Validatio
Epoch[3/50], Loss:2.300522804260254, Training Accuracy: 9.945%, Validation

```

```

# Train a dense neural network
dense_model = DenseNet(num_class=10).to(device)
criterion = nn.CrossEntropyLoss()
learning_rate = 0.1
optimizer = torch.optim.SGD(params=dense_model.parameters(), lr=learning_rate)
train_loader = train_loader
valid_loader = valid_loader
num_epoch = 30
n_no_improvement = 5
dense_model, loss_list, train_accuracy_list, valid_accuracy_list = train(dense_model,
learning_curve(save=True, path="dense_model_lc", title="DNN learning curve", train_loader=train_loader,
valid_loader=valid_loader, criterion=criterion, optimizer=optimizer, num_epoch=num_epoch, n_no_improvement=n_no_improvement))

```




```
Epoch[1/30], Loss:1.726396918296814, Training Accuracy: 30.885%, Validation
Epoch[2/30], Loss:1.938817024230957, Training Accuracy: 37.28%, Validation
Epoch[3/30], Loss:1.8413617610931396, Training Accuracy: 39.1475%, Validati
Epoch[4/30], Loss:1.6248738765716553, Training Accuracy: 42.78%, Validation
Epoch[5/30], Loss:1.7452144622802734, Training Accuracy: 41.4825%, Validati
Epoch[6/30], Loss:1.5426400899887085, Training Accuracy: 49.075%, Validatio
Epoch[7/30], Loss:1.4607850313186646, Training Accuracy: 47.557500000000005
Epoch[8/30], Loss:1.3142403364181519, Training Accuracy: 52.68000000000001%
Epoch[9/30], Loss:1.5302914381027222, Training Accuracy: 52.845%, Validatio
Epoch[10/30], Loss:1.3847599029541016, Training Accuracy: 55.75499999999999
Epoch[11/30], Loss:1.596746802330017, Training Accuracy: 53.86%, Validation
Epoch[12/30], Loss:1.2815014123916626, Training Accuracy: 56.645%, Validati
Epoch[13/30], Loss:1.1623497009277344, Training Accuracy: 58.34249999999999
Epoch[14/30], Loss:1.1158604621887207, Training Accuracy: 60.5%, Validation
Epoch[15/30], Loss:1.0876585245132446, Training Accuracy: 62.57%, Validatio
Epoch[16/30], Loss:1.1641464233398438, Training Accuracy: 61.4425%, Validat
Epoch[17/30], Loss:1.0584123134613037, Training Accuracy: 66.9775%, Validat
Epoch[18/30], Loss:1.0207831859588623, Training Accuracy: 67.6225%, Validat
Epoch[19/30], Loss:1.0777592658996582, Training Accuracy: 68.93%, Validatio
Epoch[20/30], Loss:0.8432711958885193, Training Accuracy: 65.66%, Validatio
```

DNN learning curve




```
# Train a CNN with no max pooling layers
start_time = time.time() # Calculate the execution time
cnn_model_nomax = ConvNetNoMax(num_class=10).to(device)
criterion = nn.CrossEntropyLoss()
learning_rate = 0.1
optimizer = torch.optim.SGD(params=cnn_model_nomax.parameters(), lr=learning_rate)
train_loader = train_loader
valid_loader = valid_loader
num_epoch = 30
n_no_improvement = 3
cnn_model_nomax, loss_list, train_accuracy_list, valid_accuracy_list = train(cnn_mo
print("--- %s seconds ---" % (time.time() - start_time)) # Calculate the execution
learning_curve(save=False, path="cnn_model_nomax_lc", title="CNN(without max poolin
```

```
Epoch[1/30], Loss:1.4803636074066162, Training Accuracy: 45.81%, Validation
Epoch[2/30], Loss:1.0424985885620117, Training Accuracy: 63.375%, Validation
Epoch[3/30], Loss:1.00474214553833, Training Accuracy: 68.5975%, Validation
Epoch[4/30], Loss:0.9432967305183411, Training Accuracy: 73.85249999999999%
Epoch[5/30], Loss:1.0646564960479736, Training Accuracy: 77.47749999999999%
Epoch[6/30], Loss:0.7092303037643433, Training Accuracy: 84.22500000000001%
Epoch[7/30], Loss:0.39393460750579834, Training Accuracy: 91.13250000000001
Epoch[8/30], Loss:0.4810178279876709, Training Accuracy: 91.44%, Validation
Epoch[9/30], Loss:0.31664401292800903, Training Accuracy: 92.7525%, Validat
Epoch[10/30], Loss:0.21612733602523804, Training Accuracy: 94.965%, Validat
--- 548.0471804141998 seconds ---
```



```
# Train a CNN with 2 convolutional blocks
start_time = time.time() # Calculate the execution time
cnn_model2layer = ConvNet2layer(num_class=10).to(device)
criterion = nn.CrossEntropyLoss()
learning_rate = 0.1
optimizer = torch.optim.SGD(params=cnn_model2layer.parameters(), lr=learning_rate)
train_loader = train_loader
valid_loader = valid_loader
num_epoch = 30
n_no_improvement = 3
cnn_model2layer, loss_list, train_accuracy_list, valid_accuracy_list = train(cnn_mo
print("--- %s seconds ---" % (time.time() - start_time))
learning_curve(save=False, path="cnn_model_nomax_lc", title="CNN(without max poolin
```

```
Epoch[1/30], Loss:1.5683300495147705, Training Accuracy: 47.745%, Validation
Epoch[2/30], Loss:1.3041530847549438, Training Accuracy: 52.66%, Validation
Epoch[3/30], Loss:1.1676266193389893, Training Accuracy: 66.1675%, Validati
Epoch[4/30], Loss:0.9332227110862732, Training Accuracy: 68.95750000000001%
Epoch[5/30], Loss:0.9660124778747559, Training Accuracy: 70.08%, Validation
Epoch[6/30], Loss:1.0733729600906372, Training Accuracy: 71.9425%, Validati
Epoch[7/30], Loss:0.9733961820602417, Training Accuracy: 75.78%, Validation
Epoch[8/30], Loss:0.7123333811759949, Training Accuracy: 75.605%, Validatio
Epoch[9/30], Loss:0.8251937031745911, Training Accuracy: 78.7025%, Validati
Epoch[10/30], Loss:0.8192508220672607, Training Accuracy: 77.345%, Validati
Epoch[11/30], Loss:0.7416654825210571, Training Accuracy: 76.9525%, Validat
Epoch[12/30], Loss:0.7831147313117981, Training Accuracy: 81.6325%, Validat
--- 145.25149631500244 seconds ---
```



```
# Finding the images with highest similarity(lowest Euclidean distance) to the targ
# in term of the output of second fully connected layer
```

```
fig, ax = plt.subplots(5,7,figsize=[25,20])
train_image = 0
row = 0
while train_image < 5: # 5 target images
    selected_image, selected_target = train_dataset.__getitem__(train_image) # get th
    selected_image = torch.unsqueeze(selected_image, 0) # Unsqueeze to add an axis to
    selected_image = selected_image.to(device)
    _ = cnn_model(selected_image) # Forward feed
    selected_image_flatten = cnn_model.get_layer(5) # Extract the 1x4096 vector
    print(classes[selected_target]) # Print and check the true class
```

```
cnn_model.eval()
index_nn_pq = PriorityQueue() # A priority queue to sort the similarity
with torch.no_grad():
    for index, num_image in enumerate(range(len(valid_dataset))):
        image, target = valid_dataset.__getitem__(num_image)
        image = torch.unsqueeze(image, 0)
        image = image.to(device)
        _ = cnn_model(image)
        image_flatten = cnn_model.get_layer(5)
        # Calculate the Euclidean distance between the target image and the current i
        n_n = torch.sum((selected_image_flatten - image_flatten)**2)
        # Put the Euclidean distance and image index into the pirority queue
        index_nn_pq.put((n_n, index))
```

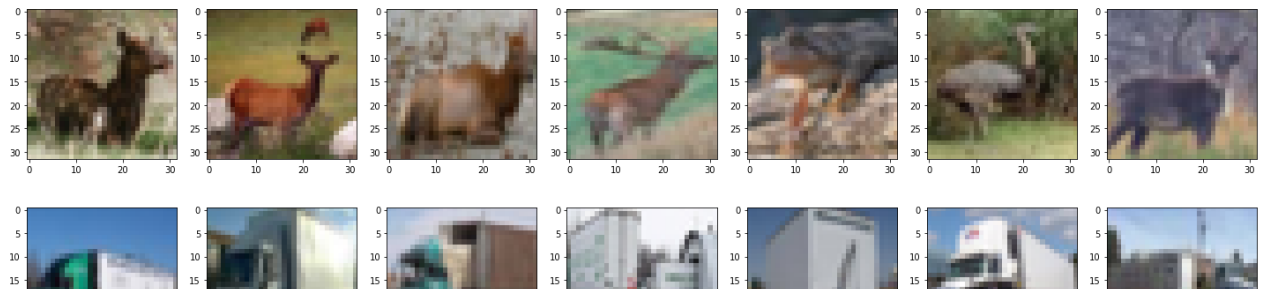
```
# Display the image
for column in range(7):
    # Display the target image
    if column == 0:
        selected_image, selected_target = train_dataset.__getitem__(train_image)
        ax[row][column].imshow(torch.permute(selected_image, (1,2,0)))
        continue
    # Display the image with the highest similarity
    n_n, index = index_nn_pq.get() # Dequeue the pirority queue
    image, target = valid_dataset.__getitem__(index) # Get the image
    ax[row][column].imshow(torch.permute(image, (1,2,0))) # Plot the image
```

```
    print(classes[target], n_n) # Check the class and the similarity(Euclidean dist  
row += 1  
train_image += 1
```

```

deer
deer tensor(63.7309, device='cuda:0')
deer tensor(65.1599, device='cuda:0')
deer tensor(69.8398, device='cuda:0')
frog tensor(72.1103, device='cuda:0')
bird tensor(72.1882, device='cuda:0')
deer tensor(72.5016, device='cuda:0')
truck
truck tensor(122.3137, device='cuda:0')
truck tensor(140.6585, device='cuda:0')
truck tensor(152.1940, device='cuda:0')
truck tensor(152.5486, device='cuda:0')
truck tensor(152.6234, device='cuda:0')
truck tensor(157.1622, device='cuda:0')
truck
car tensor(85.6006, device='cuda:0')
truck tensor(92.2293, device='cuda:0')
truck tensor(94.0501, device='cuda:0')
bird tensor(96.6377, device='cuda:0')
truck tensor(97.7937, device='cuda:0')
plane tensor(97.9414, device='cuda:0')
horse
horse tensor(159.9772, device='cuda:0')
horse tensor(168.3970, device='cuda:0')
horse tensor(176.3793, device='cuda:0')
bird tensor(178.8043, device='cuda:0')
cat tensor(182.6591, device='cuda:0')
dog tensor(184.1456, device='cuda:0')
deer
deer tensor(40.5932, device='cuda:0')
deer tensor(42.6102, device='cuda:0')
deer tensor(42.6633, device='cuda:0')
deer tensor(43.4705, device='cuda:0')
deer tensor(43.9916, device='cuda:0')
deer tensor(44.8067, device='cuda:0')

```



```

# Dimension reduction by T-SNE of all 1x4096 vectors from all images
fig, ax = plt.subplots(1,2,figsize=[25,15])
i = 0
for model in [cnn_model, dense_model]:
    # Put all 1x4096 vectors and their label from all images into a list
    flatten_vectors = []
    with torch.no_grad():
        for index, num_image in enumerate(range(len(valid_dataset))):
            image, target = valid_dataset.__getitem__(num_image)
            image = torch.unsqueeze(image, 0)
            image = image.to(device)
            _ = model(image)
            image_flatten = model.get_layer(5)
            flatten_vectors.append((image_flatten, target))

```

```
# Separate 1x4096 vectors and their labels into two list
vectors = []
target_vector = []
with torch.no_grad():
    for vector in flatten_vectors:
        flatten_image, target = vector
        flatten_image = flatten_image.detach().cpu().numpy().ravel()
        vectors.append(flatten_image)
        target_vector.append(target)
vectors = np.array(vectors)

# Perform T-SNE for all 1x4096 vectors to reduce them into 1x2 vectors
tsne_matrix = TSNE(n_components=2, learning_rate="auto", init="pca").fit_transfor

# Put T-SNE results and the labels into dataframe for plotting
tsne_matrix_df = pd.DataFrame(tsne_matrix, columns=["dim1", "dim2"])
tsne_matrix_df["label"] = target_vector
tsne_matrix_df["label"] = tsne_matrix_df["label"].apply(lambda x: classes[x]) # M
sns.scatterplot(x=tsne_matrix_df["dim1"], y=tsne_matrix_df["dim2"], hue=tsne_matr
if model == cnn_model:
    ax[i].set_title(f"TSNE Clustering(CNN)")
else:
    ax[i].set_title(f"TSNE Clustering(Dense)")
ax[i].set_xlabel("Dimension 1")
ax[i].set_ylabel("Dimension 2")
i += 1
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/manifold/_t_sne.py:986: FutureWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/manifold/_t_sne.py:986: FutureWarning,
```



```
# Overall evaluation of the testing set
cnn_model.eval()
with torch.no_grad():
    correct = 0
    total = 0
    for image, label in test_loader:
        image = image.to(device)
        label = label.to(device)

        output = cnn_model(image)
        _, predicted = torch.max(output.data, 1)
        total += label.size(0)
        correct += (predicted == label).sum().item()
    print(f"Accuracy: {correct/total*100}%")
```

Accuracy: 69.74000000000001%

```
# Evaluation function for a distribution of correct predictions in each class
def evaluation_class(model, dataset):
    correct_array = np.zeros(10) # Correc count for 10 classes
    total_array = np.zeros(10) # Total count for 10 classes
    with torch.no_grad():
        for i in range(len(test_dataset)):
            image, label = test_dataset.__getitem__(i)
            total_array[label] += 1
            image = torch.unsqueeze(image, 0)
            image = image.to(device)
            output = model(image)
            _, predicted = torch.max(output.data, 1) # Get the index of the highest score
```

```

if predicted == label: # Count correct predictions
    correct_array[predicted] += 1
return correct_array, total_array



```

```

# Plot the distribution
fig, ax = plt.subplots(1,2,figsize=[15,8])
correct_array_cnn, total_array_cnn = evaluation_class(cnn_model, test_dataset)
correct_array_dense, total_array_dense = evaluation_class(dense_model, test_dataset)
sns.barplot(x=classes, y=correct_array_cnn/total_array_cnn, ax=ax[0])
sns.barplot(x=classes, y=correct_array_dense/total_array_dense, ax=ax[1])
ax[0].set_title("Correct prediction distribution(CNN)")
ax[0].set_xlabel("Class")
ax[0].set_ylabel("Accuracy(%)")
ax[0].set_ylim([0, 1])
ax[1].set_title("Correct prediction distribution(Dense)")
ax[1].set_xlabel("Class")
ax[1].set_ylabel("Accuracy(%)")
ax[1].set_ylim([0, 1])

```

(0.0, 1.0)



```

# Plot the distribution of predicting cat images
prediction = np.zeros(10)
total = np.zeros(10)
with torch.no_grad():
    for i in range(len(test_dataset)):

```